

NoSQL & Amazon DynamoDB

DSCI 551

Wensheng Wu

Notes

- UI
 - genre: **fiction**, non-fiction
 - rating: **5**, 4, 3, 2, 1
 - **search** button
- data (books):
 - not sorted: 5, 4, 5, 3, 4, 5, ...
 - sorted by rating: 5, 5, 5, 4, 4, 3, ...

Notes (binary search)

- 5, 3, 2, 8, 9, 2, 8, 7, 5, 2, 3
- $x = 8$
- 9, 8, 8, 7, 5, 5, 3, 3, 2, 2, 2
 - midpoint: 5 \Rightarrow left (9, 8, 8, 7, 5)
 - midpoint: 8 \Rightarrow right here

Notes

- author, e.g., $h(\text{john}) = (106 + 111 + 104 + 110) \% 2 = 1$
 - host 0: storing all books with even isbn's
 - bill, bill, john, john, john, mary, ...
 - sort books by year
 - (john, 2010), (john, 2015), (john, 2020), (john, 2021), ...
 - 2010, 2015, 2020, 2021
 - host 1: storing odd isbn's
 - also sort by year
- why partition?
- why sort?

ASCII

Dec	Hex	Name	Char	Ctrl-char	Dec	Hex	Char	Dec	Hex	Char	Dec	Hex	Char
0	0	Null	NUL	CTRL-@	32	20	Space	64	40	@	96	60	`
1	1	Start of heading	SOH	CTRL-A	33	21	!	65	41	A	97	61	a
2	2	Start of text	STX	CTRL-B	34	22	"	66	42	B	98	62	b
3	3	End of text	ETX	CTRL-C	35	23	#	67	43	C	99	63	c
4	4	End of xmit	EOT	CTRL-D	36	24	\$	68	44	D	100	64	d
5	5	Enquiry	ENQ	CTRL-E	37	25	%	69	45	E	101	65	e
6	6	Acknowledge	ACK	CTRL-F	38	26	&	70	46	F	102	66	f
7	7	Bell	BEL	CTRL-G	39	27	'	71	47	G	103	67	g
8	8	Backspace	BS	CTRL-H	40	28	(72	48	H	104	68	h
9	9	Horizontal tab	HT	CTRL-I	41	29)	73	49	I	105	69	i
10	0A	Line feed	LF	CTRL-J	42	2A	*	74	4A	J	106	6A	j
11	0B	Vertical tab	VT	CTRL-K	43	2B	+	75	4B	K	107	6B	k
12	0C	Form feed	FF	CTRL-L	44	2C	,	76	4C	L	108	6C	l
13	0D	Carriage feed	CR	CTRL-M	45	2D	-	77	4D	M	109	6D	m
14	0E	Shift out	SO	CTRL-N	46	2E	.	78	4E	N	110	6E	n
15	0F	Shift in	SI	CTRL-O	47	2F	/	79	4F	O	111	6F	o
16	10	Data line escape	DLE	CTRL-P	48	30	0	80	50	P	112	70	p
17	11	Device control 1	DC1	CTRL-Q	49	31	1	81	51	Q	113	71	q
18	12	Device control 2	DC2	CTRL-R	50	32	2	82	52	R	114	72	r
19	13	Device control 3	DC3	CTRL-S	51	33	3	83	53	S	115	73	s
20	14	Device control 4	DC4	CTRL-T	52	34	4	84	54	T	116	74	t
21	15	Neg acknowledge	NAK	CTRL-U	53	35	5	85	55	U	117	75	u
22	16	Synchronous idle	SYN	CTRL-V	54	36	6	86	56	V	118	76	v
23	17	End of xmit block	ETB	CTRL-W	55	37	7	87	57	W	119	77	w
24	18	Cancel	CAN	CTRL-X	56	38	8	88	58	X	120	78	x
25	19	End of medium	EM	CTRL-Y	57	39	9	89	59	Y	121	79	y
26	1A	Substitute	SUB	CTRL-Z	58	3A	:	90	5A	Z	122	7A	z
27	1B	Escape	ESC	CTRL-[59	3B	;	91	5B	[123	7B	{
28	1C	File separator	FS	CTRL-\	60	3C	<	92	5C	\	124	7C	
29	1D	Group separator	GS	CTRL-]	61	3D	=	93	5D]	125	7D	}
30	1E	Record separator	RS	CTRL-^	62	3E	>	94	5E	^	126	7E	~
31	1F	Unit separator	US	CTRL-`	63	3F	?	95	5F	_	127	7F	DEL

your app

- isbn: _100____ (search)

Roadmap

- NoSQL 
- Amazon DynamoDB

Relational databases

- Mature & stable
 - Suitable for mission-critical applications, e.g., banking
- Feature-rich versatile query language: SQL
- ACID properties
 - In particular, strong consistency

ACID

- Atomicity: Either all or none of operations in the transaction should be executed
- Consistency: After transaction completes, the database is in a consistent state
- Isolation: allow concurrent execution of multiple transactions that do not interfere with each other
 - locking protocol
- Durability: can recover from failure
 - logging protocol

notes

- checking: 1000 saving: 1000
- transaction: transfer \$500 from checking to saving
 - step 1: deduct 500 from checking
 - checking: 500, saving: 1000
 - step 2: credit 500 to saving
 - checking: 500, saving: 1500

Strong consistency

- Traditionally, a database transaction needs to satisfy ACID properties
 - 'C' in ACID for strong consistency
- Consider a balance-transfer transaction
 - \$500 from account A to account B
 - After transfer, the total balance remains the same
 - & users do not get to see the inconsistent state (e.g., debit \$500 from A, not yet credit B)

Challenges

- Internet-scale systems & applications
 - E-commerce systems (e.g., Amazon)
 - Social media apps (e.g., Facebook, LinkedIn, Instagram)
- Big data
 - Often unstructured or semi-structured
- New workloads
 - Write/update-heavy
 - Demand high availability
 - Can tolerate weak consistency

Eventual consistency

- If no new updates are made to the object, **eventually** all accesses to the object will return the last updated value.
- A form of **weak consistency**
 - Allow users to see the inconsistency state
 - Needed to achieve high availability (HA)

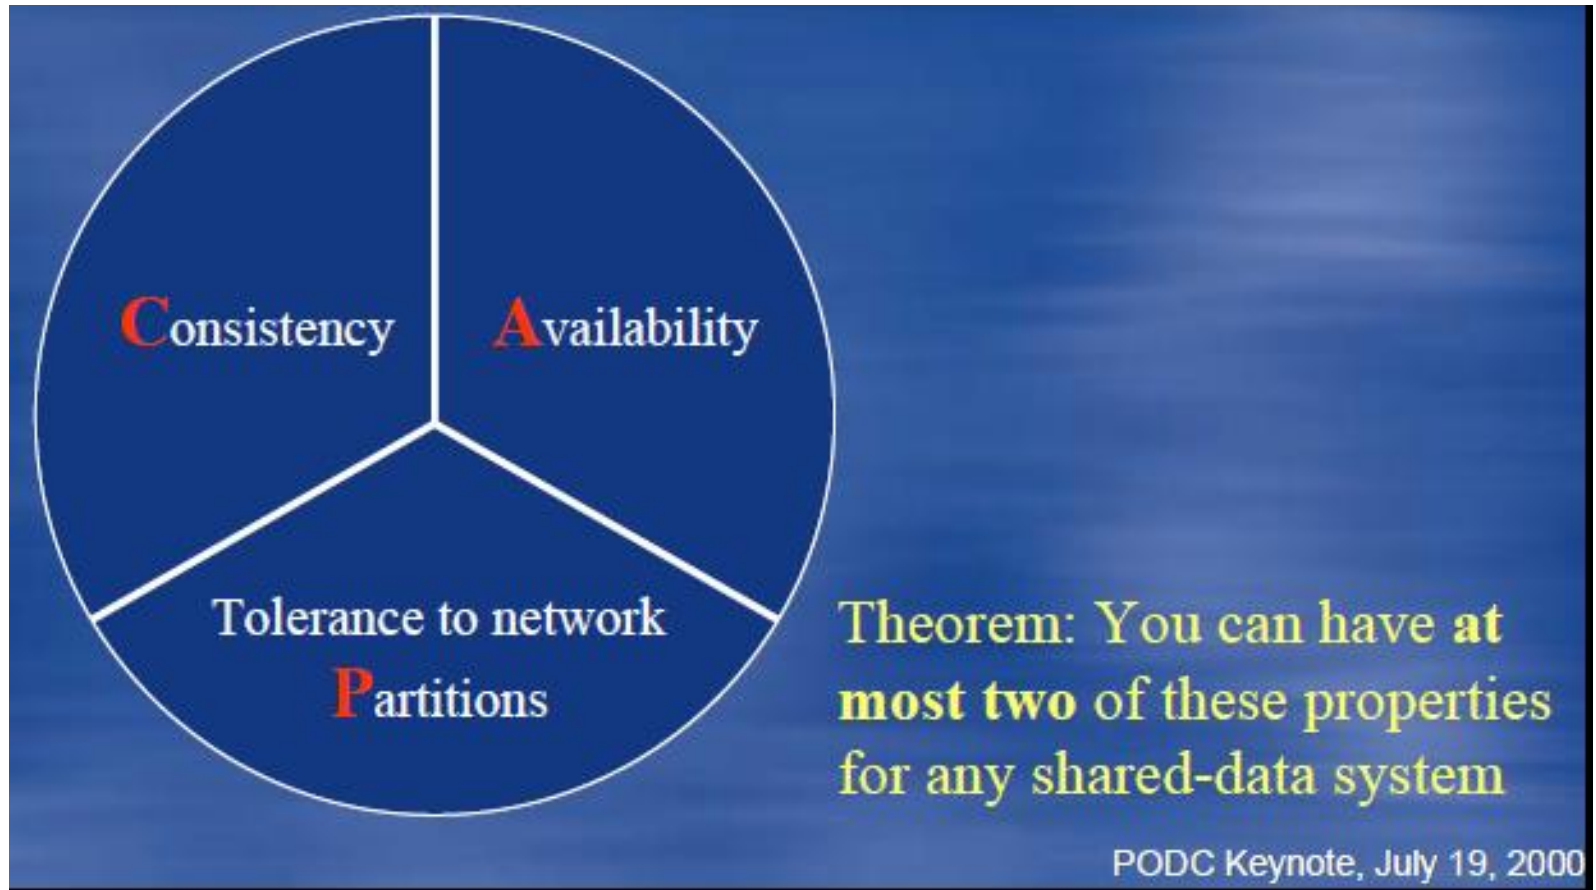
Inconsistency window

- Time between update acknowledged to user and eventual consistency achieved
 - i.e., updates propagated to all replicas
- Length of window determined by:
 - Communication delay
 - Load on the system
 - Number of replicas

Example

- DNS (domain name system) implements eventual consistency
 - E.g., DNS resolves www.usc.edu to 128.125.253.146
- Permissible for some DNS servers to have old data
 - As long as updates eventually propagated to them

CAP theorem



Explanation

Strong consistency



<i>Consistency</i>	<u><i>Availability</i></u>	<u><i>Partition tolerance</i></u>
Every read receives the most recent write or an error	Every request receives a (non-error) response – without guarantee that it contains the most recent write	The system continues to operate despite an arbitrary number of messages being dropped (or delayed) by the network between nodes

Consequence

- A distributed system needs to tolerate partitioning
 - In other words, property P is required
- Thus, when the network is partitioned, we need to choose between availability and (strong) consistency
 - ⇒ viability of eventual consistency model

Consequence

- Consider update made to an object O
- User A in LA may see the updated O right away
- But user B in NYC may see the old value of O
 - At least for a while

Eventual consistency model

- Acceptable to many applications
 - E.g., social media, cloud data storage, e-commerce
- Examples:
 - Amazon S3
 - Amazon DynamoDB (backbone of Amazon e-commerce and Web services)

NoSQL databases

- NoSQL: Not only SQL
- Key features
 - Flexible (non-relational) data model
 - Can be easily scaled out (horizontal scalability)
 - Data replicated over multiple servers
 - Weaker consistency model
 - High availability

Scale out vs. scale up

- Scale up (vertical scaling)
 - Beefing up a computer system
 - E.g., adding more CPUs, RAMs, and storage
- Scale out (horizontal scaling)
 - Adding more (commodity) computers
 - Moving some data to new computers

Types of NoSQL databases

- Key-value stores

- [Redis](#)

```
127.0.0.1:6379> set usc 'hello world'  
OK  
127.0.0.1:6379> get usc  
"hello world"
```

- Document stores

- Firebase: entire database is a JSON value
- MongoDB: database -> collections/tables -> JSON docs
- DynamoDB: database -> tables -> rows -> key-value pairs

- Wide column stores

- Database -> tables -> rows & columns
- Different rows may have different columns
- E.g., Apache Cassandra & HBase

Roadmap

- NoSQL
- Amazon DynamoDB




Amazon DynamoDB

- Schema-less: no predefined schema
 - Other than primary key
- Database contains a list of tables, e.g., music
- A table consists of a set of items/rows
 - E.g., a set of music CDs
- Each item contains a set of attributes
 - E.g., artist, title, year of CD

Items

- Similar to rows in relational databases
- But different rows may have different set of attributes
- Max size of an item: 400K
- No concept of columns in DynamoDB

DynamoDB table structure




Primary key

- Each item is uniquely identified by a primary key
- Primary key consists of
 - partition key
 - (optional) sort key

Partition key

- Partition key
 - Partition (by hashing) the data across hosts for scalability & availability
- Pick an attribute with wide range of values & evenly distributed patterns for partition key
 - E.g., user ID
- E.g., artist name
 - Hash function may put "Rod Stewart" and "Maria Kelly" in the same partition

Sort key

- Allow searching within a partition
 - E.g., year
 - So primary key = artist + year
 - This allows search all CDs by a specific artist and produced in certain years
- Possible multiple items with the same artist but different years
- 

All services

🔍 *Find services by names, keywords or acronyms.*

AWS Backup

technologies

Amazon Braket

Database

RDS

DynamoDB

ElastiCache

Neptune

Amazon QLDB

Amazon DocumentDB

Amazon Keyspaces

Management & Governance

AWS Organizations

CloudWatch

AWS Auto Scaling

CloudFormation

CloudTrail

Example

Table name*

Books



Primary key*

Partition key

Author

String



☒ Add sort key

Year

Number



Example (may vary in new version)

The screenshot displays a web interface for managing data items. At the top, a horizontal navigation bar contains tabs for 'Overview', 'Items' (which is selected and highlighted with an orange border), 'Metrics', 'Alarms', 'Capacity', 'Indexes', and a 'More' link with a downward arrow. Below the navigation bar, there is a row of controls: a blue 'Create item' button (highlighted with a red rectangle), a grey 'Actions' button with a downward arrow, and two icons (a gear and a refresh symbol) on the right. The main content area has a header bar with the text 'Scan: [Table] Books: Author, Year' and an upward arrow, followed by 'Viewing 0 to 0 items'. Below this, there is a search section with a 'Scan' dropdown menu, a text input field containing '[Table] Books: Author, Year', and a '+ Add filter' button. A 'Start search' button is located below the input field. At the bottom, a table is partially visible with columns for 'Author' and 'Year', and a checkbox in the first column.

Overview Items Metrics Alarms Capacity Indexes More ▾

Create item Actions ▾

Scan: [Table] Books: Author, Year ⬆ Viewing 0 to 0 items

Scan ▾ [Table] Books: Author, Year



+ Add filter

Start search

<input type="checkbox"/>	Author	Year
--------------------------	--------	------

Example

Create item

Tree ▾  

▼ Item {3}

+

 Author String : Jeffrey Ullman

+

 Year Number : 2005

+

 Title String : Database systems: a complete book

May add new attributes

Example


```
+ Author String : Bill Clinton
+ Year Number : 2002
+ Title String : My life
+ ▼ Keywords List [3]
+   0 String : History
+   1 String : President
+   2 Number : 2002
+ ▼ Prices NumberSet [2]
+   0 : 53.88
+   1 : 55.75
```

Value can be a list
or a set

List: ordered, heterogeneous
Set: unordered, homogeneous

Example

Data types



	Title	Keywords	Prices
	My life	[{ "S" : "History" }, { "S" : "President" }, { "N" : "2002" }]	{ 53.88, 55.75 }
Databas...			



Example

Map: contains a list of key-value pairs

Tree ▾

▼ Item {4}

- ⊕ Author String : Jeffrey Ullman
- ⊕ ▼ Ratings Map {2} ← Value can also be a map
 - ⊕ Amazon String : 5
 - ⊕ Barnes & Noble Number : 4.5map
- ⊕ Title String : Database systems: a complete book
- ⊕ Year Number : 2005

Available data types for values

String

Binary

Number

StringSet

NumberSet

BinarySet

Map


List

Boolean

Null

Query

Query: [Table] Books: Author, Year ^Viewing 1 to 1 items

Query 

[Table] Books: Author, Year

Partition key

Author

String

=


Jeffrey Ullman

Sort key


Year

Number

=



2005


 Add filter

<

>

<input type="checkbox"/>	Author	Year	Title	
<input type="checkbox"/>	Jeffrey Ullman	2005	Database syst...	


Scan


Scan 

[Table] Books: Author, Year


Filter

Year

Number 

> 

2000

 Add filter

Start search

<input type="checkbox"/>	Author	Year	Title	Keywords
<input type="checkbox"/>	Jeffrey Ullman	2005	Database syst...	
<input type="checkbox"/>	Bill Clinton	2002	My life	[{ "S" : "History" }, { "S" : "President" },

PartiQL

- Insert:
 - insert into books value {'author': 'trump1', 'year': 2021}
- Select:
 - select * from books where instock = true

PartiQL

- Update:
update books
set title = 'the art of deal' // a new attribute
where author = 'trump' and year = 2021;
- Delete:
delete from books
where author = 'trump' and year = 2021

Example

The image shows a web-based SQL query editor interface. On the left, there is a sidebar titled "Resources" with a search bar labeled "Find tables". Below the search bar, it says "Tables (1)" and lists a table named "book1". The main area on the right is titled "Query 1" and contains a SQL query: `select * from books where year between 2019 and 2020`. At the bottom of the main area, there are two buttons: "Run" and "Clear".

Resources

Find tables

Tables (1)

book1

Query 1

```
1 select * from books where year between 2019 and 2020
```

Run Clear

References

- PartiQL for DynamoDB:
 - <https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/ql-reference.html>
- Working with Tables, Items, Queries, Scans, and Indexes
 - <https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/WorkingWithDynamo.htm>
!